

Project 2: Crop Production Analysis in India

Project Title	Crop Production Analysis in India
Tools	Jupyter Notebook, Spyder, Excel
Technologies	Data Science
Domain	Agriculture

Problem Statement:

The Agriculture business domain, as a vital part of the overall supply chain, is expected to highly evolve in the upcoming years via the developments, which are taking place on the side of the Future Internet. This paper presents a novel Business-to-Business collaboration platform from the agri-food sector perspective, which aims to facilitate the collaboration of numerous stakeholders belonging to associated business domains, in an effective and flexible manner.

This dataset provides a huge amount of information on crop production in India ranging from several years. Based on the Information the ultimate goal would be to predict crop production and find important insights highlighting key indicators and metrics that influence crop production.

Tools:

- Jupyter Notebook:

Jupyter Notebook was primarily used for data cleaning, exploratory data analysis (EDA), and visualization. Its interactive environment allowed for efficient step-by-step exploration of the dataset, making it easier to identify trends, patterns, and potential data issues. The flexibility of integrating code, markdown, and visual output in the same interface helped in creating a comprehensive narrative of the analysis process.

- Spyder:

Spyder was used for model building and performance evaluation. With its powerful debugging tools and integrated development environment (IDE) for Python, Spyder enabled the implementation of machine learning algorithms and allowed for effective tuning of model parameters. It was particularly useful for writing and executing larger codebases focused on optimizing and refining the prediction models.

1) Data Cleaning:

In the data cleaning process, I began by exploring the dataset to understand its structure and addressed missing values by removing rows with missing "Production" data. I then created a new column for "Yield," calculated as the ratio of production to the area. To ensure consistency, I cleaned the crop names by removing extra spaces and categorized them into broader groups like cereals, fruits, and vegetables. This categorization helped organize the data for easier analysis, and I saved the cleaned dataset for further use.

DataCleaning.ipynb:

```
[5]: import numpy as np
import pandas as pd

[6]: df = pd.read_csv("Crop Production data.csv")

[7]: df.head()
```

	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production
0	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Arecanut	1254.0	2000.0
1	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Other Kharif pulses	2.0	1.0
2	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Rice	102.0	321.0
3	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Banana	176.0	641.0
4	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Cashewnut	720.0	165.0

```
[9]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 246091 entries, 0 to 246090
Data columns (total 7 columns):
 #   Column              Non-Null Count  Dtype  
---  --
 0   State_Name          246091 non-null object  
 1   District_Name        246091 non-null object  
 2   Crop_Year            246091 non-null int64  
 3   Season               246091 non-null object  
 4   Crop                 246091 non-null object  
 5   Area                 246091 non-null float64 
 6   Production           242361 non-null float64 
dtypes: float64(2), int64(1), object(4)
memory usage: 13.1+ MB

[10]: df.describe()
```

	Crop_Year	Area	Production
count	246091.000000	2.460910e+05	2.423610e+05
mean	2005.643018	1.200282e+04	5.825034e+05
std	4.952164	5.052340e+04	1.706581e+07
min	1997.000000	4.000000e-02	0.000000e+00
25%	2002.000000	8.000000e+01	8.800000e+01
50%	2006.000000	5.820000e+02	7.290000e+02
75%	2010.000000	4.392000e+03	7.023000e+03
max	2015.000000	8.580100e+06	1.250800e+09

```
[6]: df.isnull().sum()

State_Name      0
District_Name    0
Crop_Year        0
Season           0
Crop             0
Area             0
Production      3730
dtype: int64
```

```
[12]: df = df.dropna(subset=["Production"])
```

```
[13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 242361 entries, 0 to 246090
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   State_Name    242361 non-null  object
1   District_Name 242361 non-null  object
2   Crop_Year     242361 non-null  int64
3   Season       242361 non-null  object
4   Crop         242361 non-null  object
5   Area         242361 non-null  float64
6   Production   242361 non-null  float64
dtypes: float64(2), int64(1), object(4)
memory usage: 14.8+ MB
```

```
[14]: df['Yield'] = df['Production'] / df['Area']
```

```
[10]: print(df["Crop"].unique())
```

```
['Arecanut' 'Other Kharif pulses' 'Rice' 'Banana' 'Cashewnut' 'Coconut'
'Dry ginger' 'Sugarcane' 'Sweet potato' 'Tapioca' 'Black pepper'
'Dry chillies' 'other oilseeds' 'Turmeric' 'Maize' 'Moong(Green Gram)'
'Urad' 'Arhar/Tun' 'Groundnut' 'Sunflower' 'Bajra' 'Castor seed'
'Cotton(lint)' 'Horse-gram' 'Jowar' 'Korra' 'Ragi' 'Tobacco' 'Gram'
'Wheat' 'Masoon' 'Sesamum' 'Linseed' 'Safflower' 'Onion'
'other misc. pulses' 'Samai' 'Small millets' 'Coriander' 'Potato'
'Other Rabi pulses' 'Soyabean' 'Beans & Mutter(Vegetable)' 'Bhindi'
'Brinjal' 'Citrus Fruit' 'Cucumber' 'Grapes' 'Mango' 'Orange'
'other fibres' 'Other Fresh Fruits' 'Other Vegetables' 'Papaya'
'Pome Fruit' 'Tomato' 'Mesta' 'Cowpea(Lobia)' 'Lemon' 'Pome Granet'
'Sapota' 'Cabbage' 'Rapeseed &Mustard' 'Peas (vegetable)' 'Niger seed'
'Bottle Gourd' 'Varagu' 'Garlic' 'Ginger' 'Oilseeds total' 'Pulses total'
'Jute' 'Peas & beans (Pulses)' 'Blackgram' 'Paddy' 'Pineapple' 'Barley'
'Sannhamp' 'Khesari' 'Guar seed' 'Moth' 'Other Cereals & Millets'
'Cond-spcs other' 'Turnip' 'Carrot' 'Redish' 'Arcanut (Processed)'
'Atcanut (Raw)' 'Cashewnut Processed' 'Cashewnut Raw' 'Cardamom' 'Rubber'
'Bitter Gourd' 'Drum Stick' 'Jack Fruit' 'Snak Guard' 'Tea' 'Coffee'
'Cauliflower' 'Other Citrus Fruit' 'Water Melon' 'Total foodgrain'
'Kapas' 'Colocosia' 'Lentil' 'Bean' 'Jobster' 'Perilla' 'Rajmash Kholan'
'Ricebean (nagadal)' 'Ash Gourd' 'Beet Root' 'Lab-Lab' 'Ribed Guard'
'Yam' 'Pump Kin' 'Apple' 'Peach' 'Pear' 'Plums' 'Litchi' 'Ber'
'Other Dry Fruit' 'Jute & mesta']
```

```
[17]: df['Crop'] = df['Crop'].str.strip()
```

```
[18]: # Define a dictionary to map crops to categories
```

```
crop_categories = {
    'Arecanut': 'Nuts',
    'Other Kharif pulses': 'Pulses',
    'Rice': 'Cereals',
    'Banana': 'Fruits',
    'Cashewnut': 'Nuts',
    'Coconut': 'Nuts',
    'Dry ginger': 'Cereals'
}

# Create a new column in the dataframe for crop categories
df['Crop_Category'] = df['Crop'].map(crop_categories)

df[['Crop', 'Crop_Category']].head()
```

```
[18]:
```

	Crop	Crop_Category
0	Arecanut	Nuts
1	Other Kharif pulses	Pulses
2	Rice	Cereals
3	Banana	Fruits
4	Cashewnut	Nuts

```
[19]: print(df["Crop"].nunique())
print(df["Crop_Category"].nunique())
```

```
124
10
```

```
[20]: print(df.isnull().sum())
print(df.info())
```

```
State_Name    0
District_Name 0
Crop_Year     0
Season       0
Crop         0
Area         0
Production   0
Yield        0
Crop_Category 0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
Index: 242361 entries, 0 to 246090
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   State_Name    242361 non-null  object
1   District_Name 242361 non-null  object
2   Crop_Year     242361 non-null  int64
3   Season       242361 non-null  object
4   Crop         242361 non-null  object
5   Area         242361 non-null  float64
6   Production   242361 non-null  float64
7   Yield        242361 non-null  float64
8   Crop_Category 242361 non-null  object
dtypes: float64(3), int64(1), object(5)
memory usage: 18.5+ MB
```

[21]: df.head(10)

[21]:

	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production	Yield	Crop_Category
0	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Arecanut	1254.0	2000.0	1.594896	Nuts
1	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Other Kharif pulses	2.0	1.0	0.500000	Pulses
2	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Rice	102.0	321.0	3.147059	Cereals
3	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Banana	176.0	641.0	3.642045	Fruits
4	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Cashewnut	720.0	165.0	0.229167	Nuts
5	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Coconut	18168.0	65100000.0	3583.223250	Nuts
6	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Dry ginger	36.0	100.0	2.777778	Spices
7	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Sugarcane	1.0	2.0	2.000000	Cash Crops
8	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Sweet potato	5.0	15.0	3.000000	Vegetables
9	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Tapioca	40.0	169.0	4.225000	Vegetables

[16]: df.to_csv("Cleaned_Crop_Production.csv", index=False)

2) Exploratory data analysis (EDA):

Overall Analysis:

- ❖ State-wise Insights:
 - **Kerala** has demonstrated the **highest** overall crop **production**, highlighting its significant role in India's agricultural output.
 - **Puducherry** is notable for achieving the **highest yield** among states, reflecting efficient crop production practices in this region.
- ❖ Year-wise Highlights:
 - The year **2011** stands out with the **highest production** and **yield**, indicating a peak period for agricultural performance across various crops.
- ❖ Crop Categories:
 - **Nuts: Coconut** leads both in production and yield within the nuts category.
 - **Pulses: Pulsestotal** represents the highest production within the pulses category.
 - **Cereals: Totalfoodgrain** is the top-performing category in cereals.
 - **Fruits: Banana** tops the production charts in the fruits category.
 - **Spices: Ginger** is the leading spice crop in terms of production.
 - **Cash Crops: Sugarcane** stands out as the highest in production among cash crops.
 - **Vegetables: Tapioca** is the leading vegetable crop in terms of production.
 - **Oilseeds: Oilseedtotal** shows the highest production figures in the oilseeds category.
 - **Fiber Crops: Jute** is the top fiber crop.
 - **Unknown: Jobster** is categorized under unknown crops but shows notable figures.
- ❖ Crop Category Specific Analysis:

-Rice:

Season: Primarily produced during the **Rabi** season.

Production: Highest **throughout the year**.

State: **Chandigarh** leads in **yield**, and **Punjab** in **production**.

Year: **2014** marks the peak for both **yield** and **production**.

-Wheat:

Season: Majorly grown in the **Rabi** season.

Production: High **throughout the year**.

State: **Chandigarh** leads in **yield**, and **Punjab** in **production**.

Year: **2012** shows the highest **yield**, and **2011** the highest **production**.

-Potato:

Season: Predominantly cultivated during the **Rabi** season.

Production: Significant figures **throughout the year**.

State: **West Bengal** excels in both **yield** and **production**.

Year: **2012** records the highest **yield**, and **2004** the highest **production**.

-Coconut:

Season: Grown **throughout the year**.

State: **Puducherry** leads in **yield**, and **Kerala** in **production**.

Year: **2011** sees the peak **yield**, and **2014** the highest **production**.

-Maize:

Season: Mainly produced during the **Autumn** and **Kharif** seasons.

State: **Maharashtra** leads in **yield**, and **Telangana** in **production**.

Year: **1991** shows the **highest yield**, with a **significant decline in subsequent years**
2013 records highest **production**.

```
[1]: import pandas as pd
import seaborn as sns

import matplotlib.pyplot as plt

[2]: df = pd.read_csv("Cleaned_Crop_Production.csv")

[3]: df.describe()
```

	Crop_Year	Area	Production	Yield
count	242361.000000	2.423610e+05	2.423610e+05	242361.000000
mean	2005.625773	1.216741e+04	5.825034e+05	41.649059
std	4.958285	5.085744e+04	1.706581e+07	817.572839
min	1997.000000	1.000000e-01	0.000000e+00	0.000000
25%	2002.000000	8.700000e+01	8.800000e+01	0.513514
50%	2006.000000	6.030000e+02	7.290000e+02	1.000000
75%	2010.000000	4.545000e+03	7.023000e+03	2.355450
max	2015.000000	8.580100e+06	1.250800e+09	88000.000000


```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 242361 entries, 0 to 242360
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   State_Name      242361 non-null object  
1   District_Name   242361 non-null object  
2   Crop_Year       242361 non-null int64  
3   Season          242361 non-null object  
4   Crop            242361 non-null object  
5   Area            242361 non-null float64 
6   Production      242361 non-null float64 
7   Yield           242361 non-null float64 
8   Crop_Category   242361 non-null object  
dtypes: float64(3), int64(1), object(5)
memory usage: 16.6+ MB
```


```
[84]: df['Crop_Category'].value_counts()
```

```
[84]: Crop_Category
Cereals      63487
Pulses       60272
Oilseeds     45752
Vegetables   25460
Spices       18404
Cash Crops   10561
Fibre Crops  7602
Fruits       6154
Nuts         4660
Unknown      9
Name: count, dtype: int64
```


```
[13]: sns.barplot(x="State_Name", y="Production", data=df,palette="hls", hue="State_Name", legend= False)
plt.xticks(rotation=90)
plt.title('Production by State Name')
plt.xlabel('State Name')
plt.ylabel('Production')
plt.show()
```



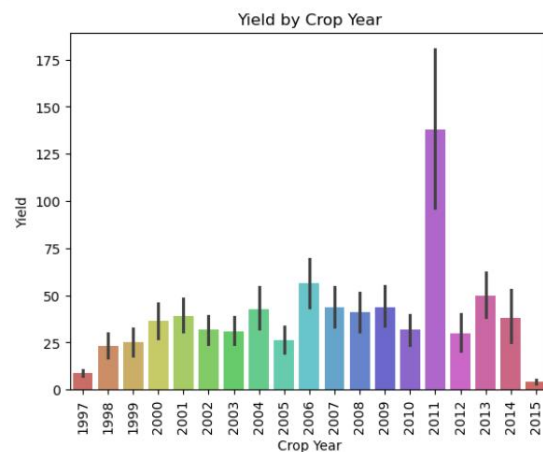
```
[14]: sns.barplot(x="State_Name", y="Yield", data=df,palette="hls", hue="State_Name", legend= False)
plt.xticks(rotation=90)
plt.title('Yield by State Name')
plt.xlabel('State Name')
plt.ylabel('Yield')
plt.show()
```



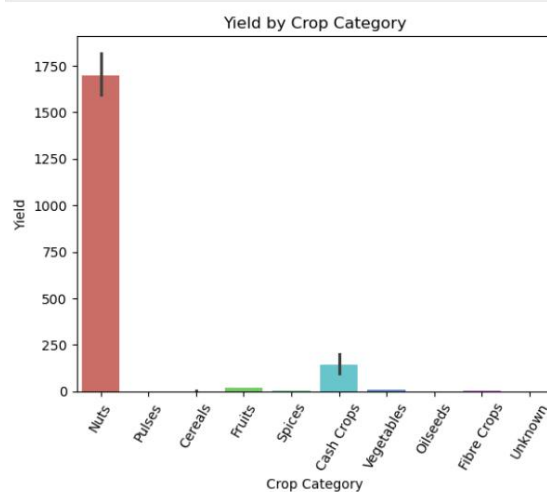
```
[15]: sns.barplot(x="Crop_Year", y="Production", data=df,palette="hls", hue="Crop_Year", legend= False)
plt.xticks(rotation=90)
plt.title('Production by Crop_Year')
plt.xlabel('Crop_Year')
plt.ylabel('Production')
plt.show()
```



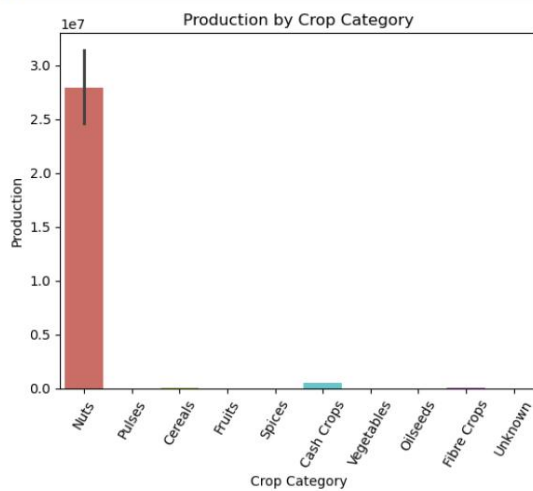
```
[16]: sns.barplot(x="Crop_Year", y="Yield", data=df,palette="hls", hue="Crop_Year", legend= False)
plt.xticks(rotation=90)
plt.title('Yield by Crop Year')
plt.xlabel('Crop Year')
plt.ylabel('Yield')
plt.show()
```



```
[17]: sns.barplot(x="Crop_Category", y="Yield", data=df,palette="hls", hue="Crop_Category", legend= False)
plt.xticks(rotation=60)
plt.title('Yield by Crop Category')
plt.xlabel('Crop Category')
plt.ylabel('Yield')
plt.show()
```



```
[18]: sns.barplot(x="Crop_Category", y="Production", data=df,palette="hls", hue="Crop_Category", legend= False)
plt.xticks(rotation=60)
plt.title('Production by Crop Category')
plt.xlabel('Crop Category')
plt.ylabel('Production')
plt.show()
```



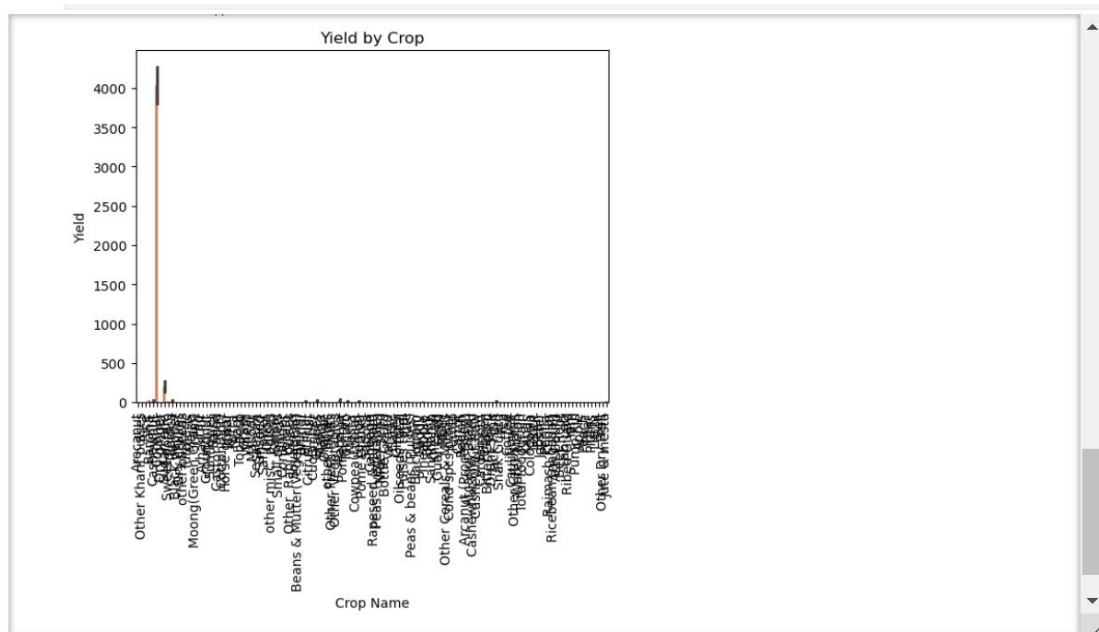
```
[91]: crop_yield = df.groupby('Crop', as_index=False)['Yield'].sum()
crop_yield['Yield'] = crop_yield['Yield'].astype(int)
sorted_crop_yield = crop_yield.sort_values(by='Yield', ascending=False)
pd.set_option('display.max_rows', None)
pd.set_option('display.float_format', '{:.0f}'.format)
print(sorted_crop_yield.to_string(index=False))

sns.barplot(x="Crop", y="Yield", data=df,palette="hls", hue="Crop", legend= False)
plt.xticks(rotation=90)
plt.title('Yield by Crop')
plt.xlabel('Crop Name')
plt.ylabel('Yield')
plt.show()
```

```

Crop  Yield
Coconut  7911110
Sugarcane  1526291
Potato  88962
Banana  84793
Onion  82797
Sweet potato  36471
Maize  36297
Rice  30088
Tapioca  29737
Wheat  16462
Dry ginger  16038
Papaya  13842
Jute  13314
Turmeric  11486
Garlic  10707
Groundnut  10480
Cashewnut  10333
Mesta  9853
Cotton(lint)  9041
Dry chillies  8812
Barley  7502
Jowar  7329
Soya  6314

```





```
[19]: crop_production = df.groupby('Crop', as_index=False)['Production'].sum()
sorted_crop_production = crop_production.sort_values(by='Production', ascending=False)
print(sorted_crop_production[['Crop', 'Production']])

sns.barplot(x="Crop", y="Production", data=df, palette="hls", hue="Crop", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Crop')
plt.xlabel('Crop Name')
plt.ylabel('Yield')
plt.show()
```

	Crop	Production
28	Coconut	1.299816e+11
106	Sugar cane	5.535682e+09
95	Rice	1.605470e+09
119	Wheat	1.332826e+09
87	Potato	4.248263e+08
..
71	Other Citrus Fruit	0.000000e+00
35	Cucumber	0.000000e+00
58	Litchi	0.000000e+00
54	Lab-Lab	0.000000e+00
0	Apple	0.000000e+00

[124 rows x 2 columns]






```
[96]: crop_categories = df['Crop_Category'].unique()



# Plot for each crop category
for category in crop_categories:
    plt.figure(figsize=(12, 8))
    category_data = df[df['Crop_Category'] == category]


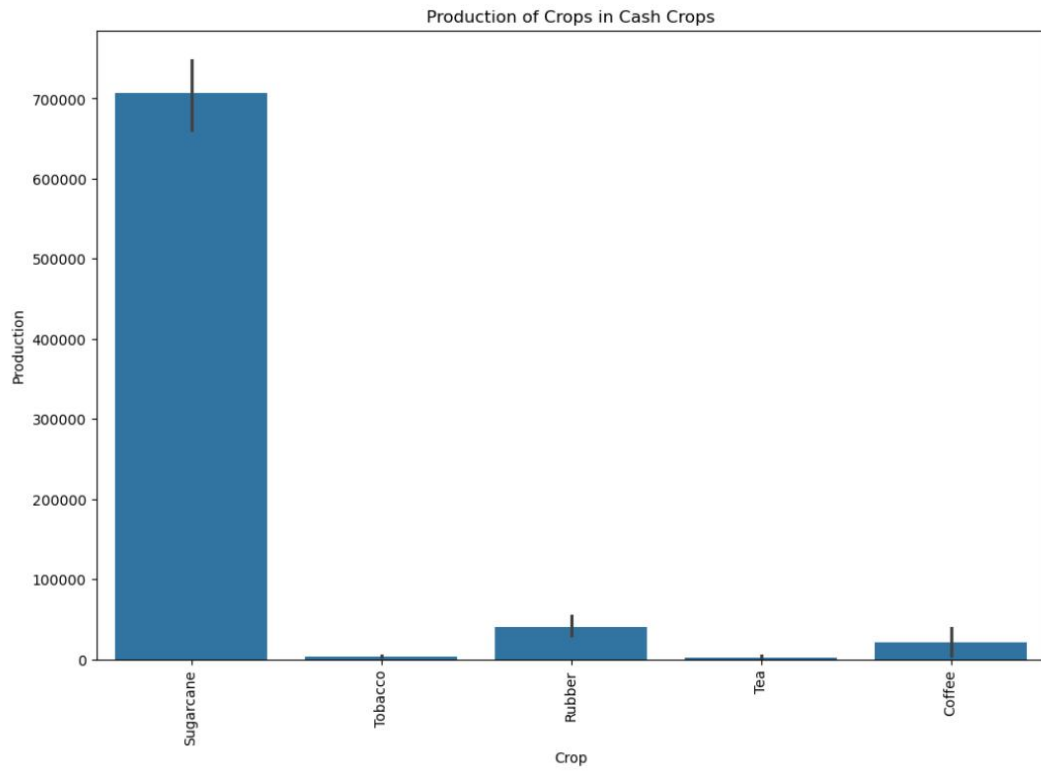
    # Optional: aggregate production if needed
    # e.g., to show average production per crop
    # category_data = category_data.groupby('Crop').agg({'Production': 'mean'}).reset_index()



    sns.barplot(x='Crop', y='Production', data=category_data)
    plt.title(f'Production of Crops in {category}')
    plt.xlabel('Crop')
    plt.ylabel('Production')
    plt.xticks(rotation=90) # Rotate x labels for better readability
    plt.show()
```

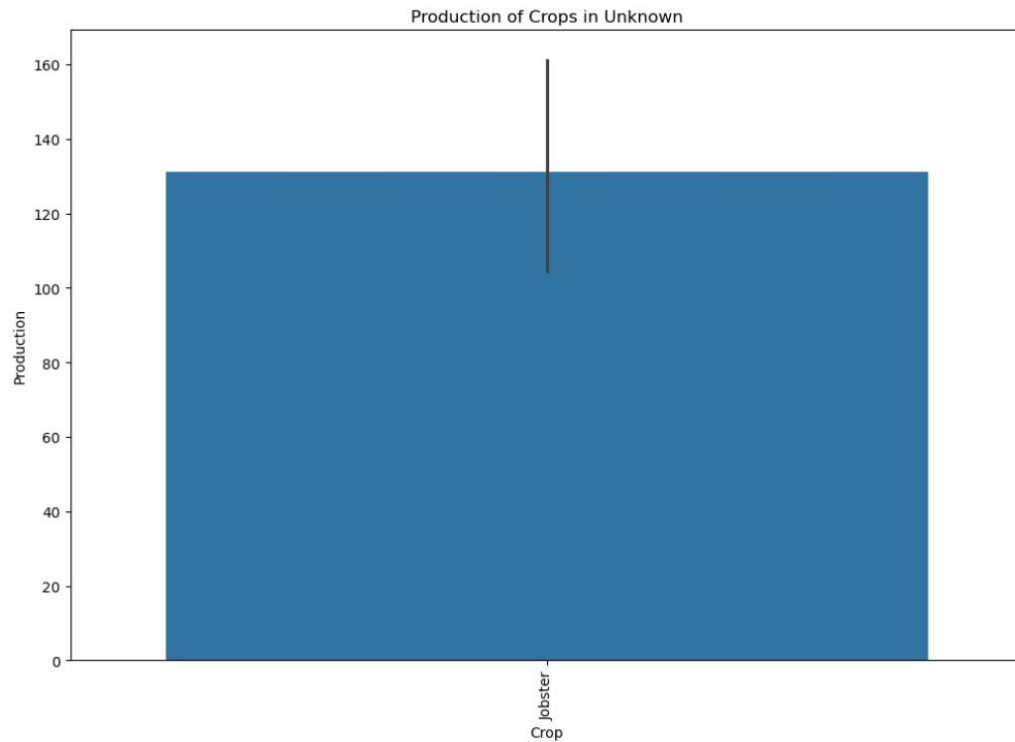












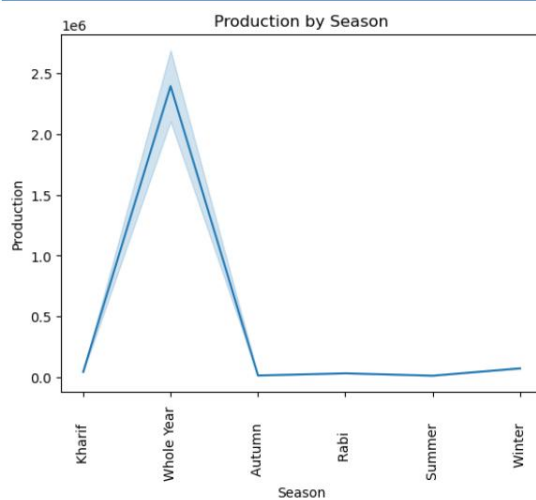
```
[21]: crop_categories = df['Crop_Category'].unique()

# Plot for each crop category
for category in crop_categories:
    plt.figure(figsize=(12, 8))
    category_data = df[df['Crop_Category'] == category]

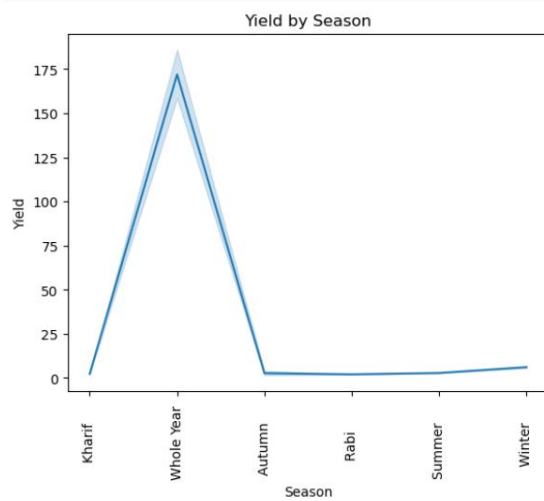
    # Optional: aggregate production if needed
    # e.g., to show average production per crop
    # category_data = category_data.groupby('Crop').agg({'Production': 'mean'}).reset_index()

    sns.barplot(x='Crop', y='Yield', data=category_data)
    plt.title(f'Production of Crops in {category}')
    plt.xlabel('Crop')
    plt.ylabel('Production')
    plt.xticks(rotation=90) # Rotate x labels for better readability
    plt.show()
```

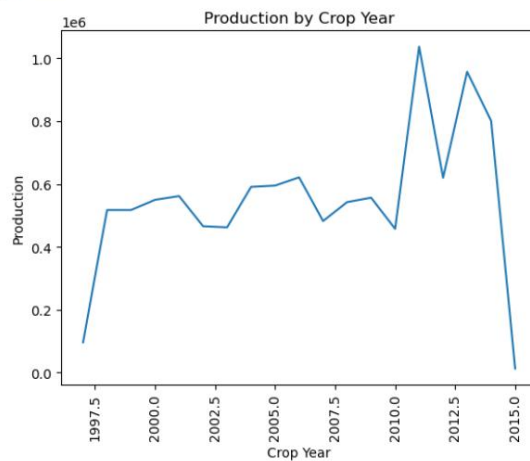
```
[23]: sns.lineplot(x="Season", y="Production", data=df)
plt.xticks(rotation=90)
plt.title('Production by Season')
plt.xlabel('Season')
plt.ylabel('Production')
plt.show()
```



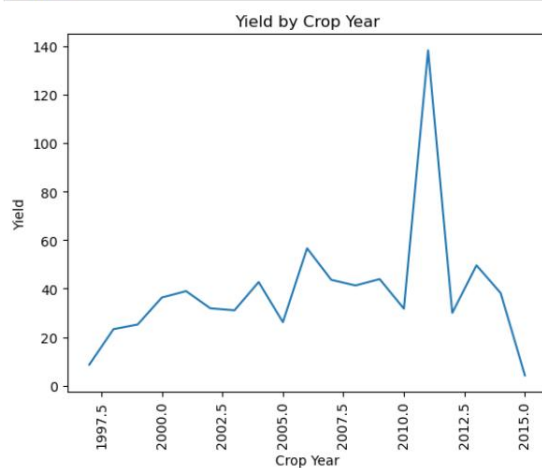
```
[24]: sns.lineplot(x="Season", y="Yield", data=df)
plt.xticks(rotation=90)
plt.title('Yield by Season')
plt.xlabel('Season')
plt.ylabel('Yield')
plt.show()
```



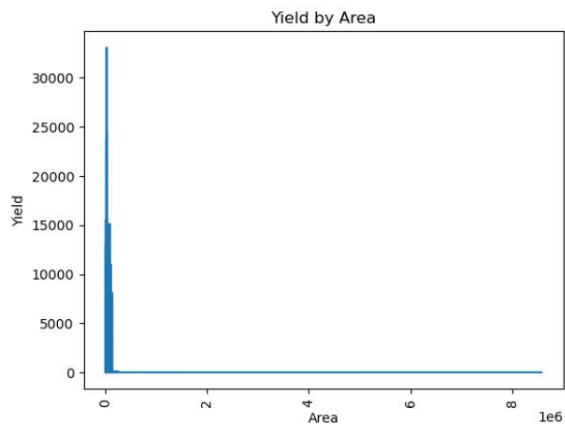
```
[25]: sns.lineplot(x="Crop_Year", y="Production", data=df, errorbar=None)
plt.xticks(rotation=90)
plt.title('Production by Crop Year')
plt.xlabel('Crop Year')
plt.ylabel('Production')
plt.show()
```



```
[26]: sns.lineplot(x="Crop_Year", y="Yield", data=df, errorbar=None)
plt.xticks(rotation=90)
plt.title('Yield by Crop Year')
plt.xlabel('Crop Year')
plt.ylabel('Yield')
plt.show()
```

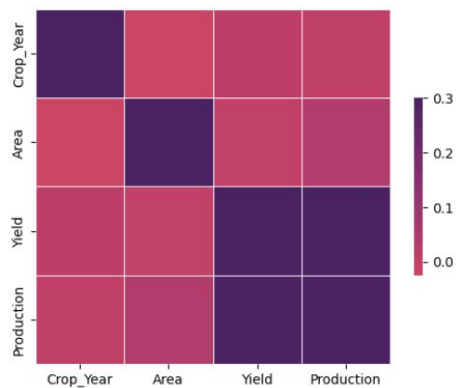



```
[27]: sns.lineplot(x="Area", y="Yield", data=df, errorbar=None)
plt.xticks(rotation=90)
plt.title('Yield by Area')
plt.xlabel('Area')
plt.ylabel('Yield')
plt.show()
```



```
[28]: cmap = sns.color_palette("flare", as_cmap=True)
sns.heatmap(df[['Crop_Year', 'Area', 'Yield', 'Production']].corr(),
            cmap=cmap,
            vmax=.3,
            center=0,
            square=True,
            linewidths=.5,
            cbar_kws={"shrink": .5})

plt.show()
```



Now analyzing Each crop (Rice, Wheat, Potato, Cotton, Maize)

1) Rice

```
[29]: df_rice_data = df[df["Crop"]=="Rice"]
df_rice_data.head()
```

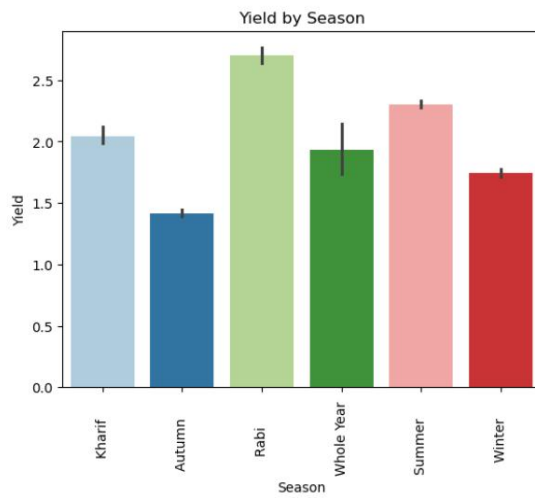
```
[29]:
```

	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production	Yield	Crop_Category
2	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Rice	102.00	321.00	3.147059	Cereals
12	Andaman and Nicobar Islands	NICOBARS	2001	Kharif	Rice	83.00	300.00	3.614458	Cereals
18	Andaman and Nicobar Islands	NICOBARS	2002	Kharif	Rice	189.20	510.84	2.700000	Cereals
27	Andaman and Nicobar Islands	NICOBARS	2003	Kharif	Rice	52.00	90.17	1.734038	Cereals
36	Andaman and Nicobar Islands	NICOBARS	2004	Kharif	Rice	52.94	72.57	1.370797	Cereals

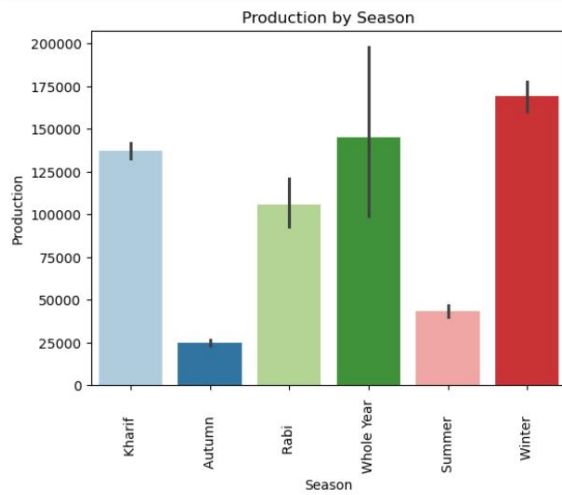
```
[30]: df_rice_data.shape
```

```
[30]: (15082, 9)
```


```
[33]: sns.barplot(x="Season", y="Yield", data=df_rice_data,palette="Paired", hue="Season", legend= False)
plt.xticks(rotation=90)
plt.title('Yield by Season')
plt.xlabel('Season')
plt.ylabel('Yield')
plt.show()
```




```
[34]: sns.barplot(x="Season", y="Production", data=df_rice_data,palette="Paired", hue="Season", legend= False)
plt.xticks(rotation=90)
plt.title('Production by Season')
plt.xlabel('Season')
plt.ylabel('Production')
plt.show()
```




```
[35]: sns.barplot(x="State_Name", y="Yield", data=df_rice_data,palette="Paired", hue="State_Name", legend= False)
plt.xticks(rotation=90)
plt.title('Yield by State Name')
plt.xlabel('State Name')
plt.ylabel('Yield')
plt.show()
```



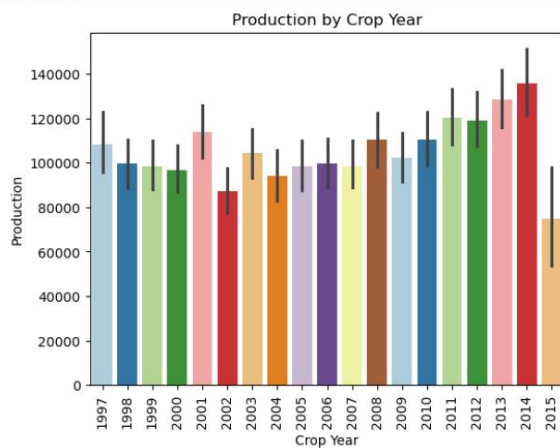
```
[36]: sns.barplot(x="State_Name", y="Production", data=df_rice_data,palette="Paired", hue="State_Name", legend= False)
plt.xticks(rotation=90)
plt.title('Production by State Name')
plt.xlabel('State Name')
plt.ylabel('Production')
plt.show()
```



```
[37]: sns.barplot(x="Crop_Year", y="Yield", data=df_rice_data,palette="Paired", hue="Crop_Year", legend= False)
plt.xticks(rotation=90)
plt.title('Yield by Crop Year')
plt.xlabel('Crop Year')
plt.ylabel('Yield')
plt.show()
```



```
[38]: sns.barplot(x="Crop_Year", y="Production", data=df_rice_data,palette="Paired", hue="Crop_Year", legend= False)
plt.xticks(rotation=90)
plt.title('Production by Crop Year')
plt.xlabel('Crop Year')
plt.ylabel('Production')
plt.show()
```



2) Wheat

```
[39]: df_wheat_data = df[df["Crop"] == "Wheat"]
df_wheat_data.head()
```

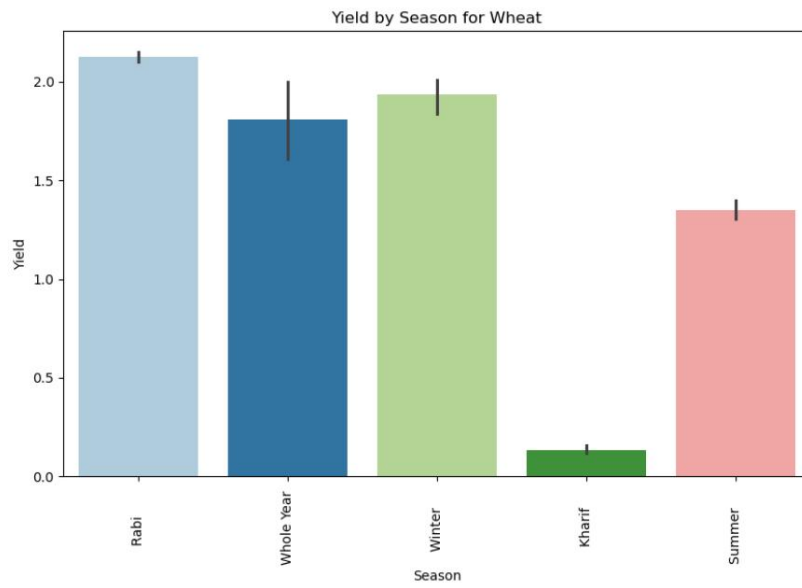
```
[39]:
```

	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production	Yield	Crop_Category
228	Andhra Pradesh	ANANTAPUR	1997	Rabi	Wheat	300.0	200.0	0.666667	Cereals
253	Andhra Pradesh	ANANTAPUR	1998	Rabi	Wheat	400.0	200.0	0.500000	Cereals
282	Andhra Pradesh	ANANTAPUR	1999	Rabi	Wheat	439.0	294.0	0.669704	Cereals
324	Andhra Pradesh	ANANTAPUR	2000	Rabi	Wheat	520.0	297.0	0.571154	Cereals
370	Andhra Pradesh	ANANTAPUR	2001	Rabi	Wheat	307.0	213.0	0.693811	Cereals

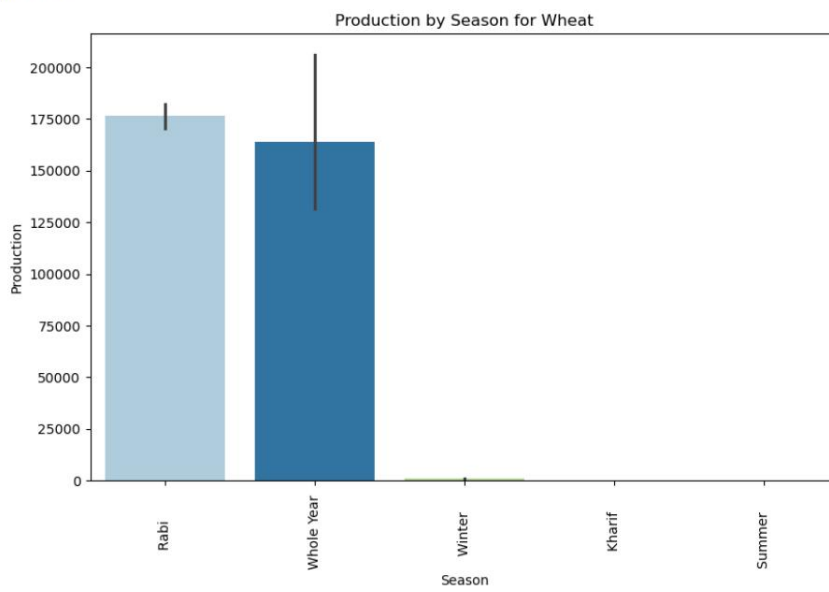
```
[40]: df_wheat_data.shape
```

```
[40]: (7878, 9)
```


```
[41]: plt.figure(figsize=(10, 6))
sns.barplot(x="Season", y="Yield", data=df_wheat_data, palette="Paired", hue="Season", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Season for Wheat')
plt.xlabel('Season')
plt.ylabel('Yield')
plt.show()
```




```
[42]: plt.figure(figsize=(10, 6))
sns.barplot(x="Season", y="Production", data=df_wheat_data, palette="Paired", hue="Season", legend=False)
plt.xticks(rotation=90)
plt.title('Production by Season for Wheat')
plt.xlabel('Season')
plt.ylabel('Production')
plt.show()
```




```
[43]: plt.figure(figsize=(10, 6))
sns.barplot(x="State_Name", y="Yield", data=df_wheat_data, palette="Paired", hue="State_Name", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by State for Wheat')
plt.xlabel('State')
plt.ylabel('Yield')
plt.show()
```




```
[44]: plt.figure(figsize=(10, 6))
sns.barplot(x="State_Name", y="Production", data=df_wheat_data, palette="Paired", hue="State_Name", legend=False)
plt.xticks(rotation=90)
plt.title('Production by State for Wheat')
plt.xlabel('State')
plt.ylabel('Production')
plt.show()
```



```
[45]: plt.figure(figsize=(10, 6))
sns.barplot(x="Crop_Year", y="Yield", data=df_wheat_data, palette="Paired", hue="Crop_Year", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Crop Year for Wheat')
plt.xlabel('Crop Year')
plt.ylabel('Yield')
plt.show()
```



```
[46]: plt.figure(figsize=(10, 6))
sns.barplot(x="Crop_Year", y="Production", data=df_wheat_data, palette="Paired", hue="Crop_Year", legend=False)
plt.xticks(rotation=90)
plt.title('Production by Crop Year for Wheat')
plt.xlabel('Crop Year')
plt.ylabel('Production')
plt.show()
```



3) Potato

```
[47]: df_potato_data = df[df["Crop"] == "Potato"]
df_potato_data.head()
```


```
[47]:
```

	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production	Yield	Crop_Category
329	Andhra Pradesh	ANANTAPUR	2000	Whole Year	Potato	4.0	34.0	8.500000	Vegetables
431	Andhra Pradesh	ANANTAPUR	2002	Whole Year	Potato	2.0	17.0	8.500000	Vegetables
528	Andhra Pradesh	ANANTAPUR	2004	Whole Year	Potato	2.0	20.0	10.000000	Vegetables
739	Andhra Pradesh	ANANTAPUR	2010	Whole Year	Potato	21.0	236.0	11.238095	Vegetables
786	Andhra Pradesh	ANANTAPUR	2011	Whole Year	Potato	18.0	181.0	10.055556	Vegetables


```
[48]: df_potato_data.shape
```

```
[48]: (6914, 9)
```


```
[49]: plt.figure(figsize=(10, 6))
sns.barplot(x="Season", y="Yield", data=df_potato_data, palette="Paired", hue="Season", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Season for Potato')
plt.xlabel('Season')
plt.ylabel('Yield')
plt.show()
```




```
[50]: plt.figure(figsize=(10, 6))
sns.barplot(x="Season", y="Production", data=df_potato_data, palette="Paired", hue="Season", legend=False)
plt.xticks(rotation=90)
plt.title('Production by Season for Potato')
plt.xlabel('Season')
plt.ylabel('Production')
plt.show()
```





```
[51]: plt.figure(figsize=(10, 6))
sns.barplot(x="State_Name", y="Yield", data=df_potato_data, palette="Paired", hue="State_Name", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by State for Potato')
plt.xlabel('State')
plt.ylabel('Yield')
plt.show()
```




```
[52]: plt.figure(figsize=(10, 6))
sns.barplot(x="State_Name", y="Production", data=df_potato_data, palette="Paired", hue="State_Name", legend=False)
plt.xticks(rotation=90)
plt.title('Production by State for Potato')
plt.xlabel('State')
plt.ylabel('Production')
plt.show()
```



```
[53]: # Yield by Crop Year
plt.figure(figsize=(18, 6))
sns.barplot(x="Crop_Year", y="Yield", data=df_potato_data, palette="Paired", hue="Crop_Year", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Crop Year for Potato')
plt.xlabel('Crop Year')
plt.ylabel('Yield')
plt.show()
```



```
[54]: plt.figure(figsize=(18, 6))
sns.barplot(x="Crop_Year", y="Production", data=df_potato_data, palette="Paired", hue="Crop_Year", legend=False)
plt.xticks(rotation=90)
plt.title('Production by Crop Year for Potato')
plt.xlabel('Crop Year')
plt.ylabel('Production')
plt.show()
```



4) Coconut

```
[55]: df_coconut_data = df[df["Crop"] == "Coconut"]
df_coconut_data
```

```
[56]:
```


	State Name	District Name	Crop Year	Season	Crop	Area	Production	Yield	Crop Category
5	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Coconut	18168.00	65100000.0	3583.233250	Nuts
14	Andaman and Nicobar Islands	NICOBARS	2001	Whole Year	Coconut	18190.00	64430000.0	3542.056075	Nuts
23	Andaman and Nicobar Islands	NICOBARS	2002	Whole Year	Coconut	18240.00	67490000.0	3700.109649	Nuts
32	Andaman and Nicobar Islands	NICOBARS	2003	Whole Year	Coconut	18284.74	68580000.0	3750.668590	Nuts
41	Andaman and Nicobar Islands	NICOBARS	2004	Whole Year	Coconut	18394.70	52380000.0	2847.559351	Nuts
...
241990	West Bengal	PURULIA	2004	Whole Year	Coconut	66.00	296.1	4.486364	Nuts
242027	West Bengal	PURULIA	2005	Whole Year	Coconut	74.00	311.0	4.202703	Nuts
242063	West Bengal	PURULIA	2006	Whole Year	Coconut	73.00	365000.0	5000.000000	Nuts
242108	West Bengal	PURULIA	2007	Whole Year	Coconut	58.00	898000.0	15482.758621	Nuts
242149	West Bengal	PURULIA	2008	Whole Year	Coconut	58.00	598.0	10.310345	Nuts

1958 rows x 9 columns


```
[56]: df_coconut_data.shape
```

```
[56]: (1958, 9)
```


```
[57]: plt.figure(figsize=(10, 6))
sns.barplot(x="Season", y="Yield", data=df_coconut_data, palette="Paired", hue="Season", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Season for Coconut')
plt.xlabel('Season')
plt.ylabel('Yield')
plt.show()
```




```
[58]: plt.figure(figsize=(10, 6))
sns.barplot(x="Season", y="Production", data=df_coconut_data, palette="Paired", hue="Season", legend=False)
plt.xticks(rotation=90)
plt.title('Production by Season for Coconut')
plt.xlabel('Season')
plt.ylabel('Production')
plt.show()
```




```
[59]: plt.figure(figsize=(10, 6))
sns.barplot(x="State Name", y="Yield", data=df_coconut_data, palette="Paired", hue="State Name", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by State for Coconut')
plt.xlabel('State')
plt.ylabel('Yield')
plt.show()
```




```
[62]: plt.figure(figsize=(10, 6))
sns.barplot(x="State_Name", y="Production", data=df_coconut_data, palette="Paired", hue="State_Name", legend=False)
plt.xticks(rotation=90)
plt.title('Production by State for Coconut')
plt.xlabel('State')
plt.ylabel('Production')
plt.show()
```



```
[63]: plt.figure(figsize=(10, 6))
sns.barplot(x="Crop_Year", y="Yield", data=df_coconut_data, palette="Paired", hue="Crop_Year", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Crop Year for Coconut')
plt.xlabel('Crop Year')
plt.ylabel('Yield')
plt.show()
```



```
[55]: plt.figure(figsize=(10, 6))
sns.barplot(x="Crop Year", y="Production", data=df_coconut_data, palette="Paired", hue="Crop Year", legend=False)
plt.xticks(rotation=90)
plt.title('Production by Crop Year for Coconut')
plt.xlabel('Crop Year')
plt.ylabel('Production')
plt.show()
```



5) Maize

```
[56]: df_maize_data = df[df["Crop"] == "Maize"]
df_maize_data.head()
```


```
[56]:
```

	State Name	District Name	Crop Year	Season	Crop	Area	Production	Yield	Crop Category
69	Andaman and Nicobar Islands	NICOBARS	2010	Rabi	Maize	3.84	18.22	4.744792	Cereals
118	Andaman and Nicobar Islands	NORTH AND MIDDLE ANDAMAN	2010	Rabi	Maize	86.70	96.40	1.111880	Cereals
192	Andaman and Nicobar Islands	SOUTH ANDAMANS	2010	Rabi	Maize	73.00	253.00	3.465753	Cereals
210	Andhra Pradesh	ANANTAPUR	1997	Kharif	Maize	2800.00	4900.00	1.750000	Cereals
224	Andhra Pradesh	ANANTAPUR	1997	Rabi	Maize	600.00	2400.00	4.000000	Cereals


```
[57]: df_maize_data.shape
```

```
[57]: (13787, 9)
```


```
[58]: plt.figure(figsize=(10, 6))
sns.barplot(x="Season", y="Yield", data=df_maize_data, palette="Paired", hue="Season", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Season for Maize')
plt.xlabel('Season')
plt.ylabel('Yield')
plt.show()
```




```
[71]: plt.figure(figsize=(10, 6))
sns.barplot(x="Season", y="Production", data=df_maize_data, palette="Paired", hue="Season", legend=False)
plt.xticks(rotation=90)
plt.title('Production by Season for Maize')
plt.xlabel('Season')
plt.ylabel('Production')
plt.show()
```




```
[72]: plt.figure(figsize=(10, 6))
sns.barplot(x="State Name", y="Yield", data=df_maize_data, palette="Paired", hue="State Name", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by State for Maize')
plt.xlabel('State')
plt.ylabel('Yield')
plt.show()
```




```
[73]: plt.figure(figsize=(18, 8))
sns.barplot(x="State_Name", y="Production", data=df_maize_data, palette="Paired", hue="State_Name", legend=False)
plt.xticks(rotation=90)
plt.title('Production by State for Maize')
plt.xlabel('State')
plt.ylabel('Production')
plt.show()
```



```
[74]: plt.figure(figsize=(18, 8))
sns.barplot(x="Crop_Year", y="Yield", data=df_maize_data, palette="Paired", hue="Crop_Year", legend=False)
plt.xticks(rotation=90)
plt.title('Yield by Crop Year for Maize')
plt.xlabel('Crop Year')
plt.ylabel('Yield')
plt.show()
```




```
[76]: plt.figure(figsize=(18, 6))
sns.barplot(x="Crop_Year", y="Production", data=df_maize_data, palette="Paired", hue="Crop_Year", legend=False)
plt.xticks(rotation=90)
plt.title('Production by Crop Year for Maize')
plt.xlabel('Crop Year')
plt.ylabel('Production')
plt.show()
```



3) Model Building:

Implemented the XGBoost algorithm to predict crop production. After splitting the data into training and testing sets, I trained the model using optimal hyperparameters with both L1 (Lasso) and L2 (Ridge) regularization to prevent overfitting. The model's performance was evaluated using metrics like R^2 and Mean Absolute Error (MAE) on both training and test datasets.

Model.py:

```
# -*- coding: utf-8 -*-
"""
```

Created on Tue Sep 3 23:18:44 2024

```
@author: ranea
"""
```

```
from sklearn.model_selection import train_test_split
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import xgboost as xgb
from sklearn.metrics import r2_score, mean_absolute_error
from sklearn.model_selection import learning_curve
```

```
df = pd.read_csv('Cleaned_Crop_Production.csv')
```

```
data = df.drop(['State_Name'], axis=1)
dummy = pd.get_dummies(data)
```

```
x = dummy.drop(["Production", "Yield"], axis=1)
```

```

y = dummy["Production"]
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=5)

print("x_train :", x_train.shape)
print("x_test :", x_test.shape)
print("y_train :", y_train.shape)
print("y_test :", y_test.shape)

xgb_model = xgb.XGBRegressor(
    reg_alpha=10,
    reg_lambda=0.1,
    n_estimators=100,
    objective='reg:squarederror',
    eval_metric='mae'
)

xgb_model.fit(x_train, y_train)

y_train_pred = xgb_model.predict(x_train)
y_test_pred = xgb_model.predict(x_test)

train_r2 = r2_score(y_train, y_train_pred)
train_mae = mean_absolute_error(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)
test_mae = mean_absolute_error(y_test, y_test_pred)

print("XGBoost Training R²:", train_r2)
print("XGBoost Training MAE:", train_mae)
print("XGBoost Test R²:", test_r2)
print("XGBoost Test MAE:", test_mae)

plt.figure(figsize=(10,6))
plt.scatter(y_test, y_test_pred, color='blue', alpha=0.6)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linestyle='--', lw=2) # 45-degree line
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('XGBoost: Test Set Actual vs Predicted Values')
plt.grid(True)
plt.show()

train_sizes, train_scores, test_scores = learning_curve(
    xgb_model, x, y, train_sizes=[100, 500, 1000, 5000, 10000], cv=5,
    scoring='neg_mean_absolute_error'
)
train_errors_mean = -train_scores.mean(axis=1)
test_errors_mean = -test_scores.mean(axis=1)

plt.figure(figsize=(10,6))
plt.plot(train_sizes, train_errors_mean, label='Training Error')
plt.plot(train_sizes, test_errors_mean, label='Test Error')
plt.xlabel('Training Set Size')
plt.ylabel('Mean Absolute Error')
plt.title('Learning Curves')
plt.legend()
plt.grid(True)
plt.show()

```

Model Result:

```
In [1]: runfile('C:/Users/ranea/CropProduction/Model.py', wdir='C:/Users/ranea/CropProduction')
x_train : (181770, 788)
x_test  : (60591, 788)
y_train : (181770,)
y_test  : (60591,)
XGBoost Training R2: 0.9965063498188226
XGBoost Training MAE: 56986.05024930907
XGBoost Test R2: 0.9599334528080093
XGBoost Test MAE: 141306.73901528507
```

